

# Metadata Load Balancing Policies and Key-Value Stores

MICHAEL SEVILLA\*

## ABSTRACT

Enter the text of your abstract here.

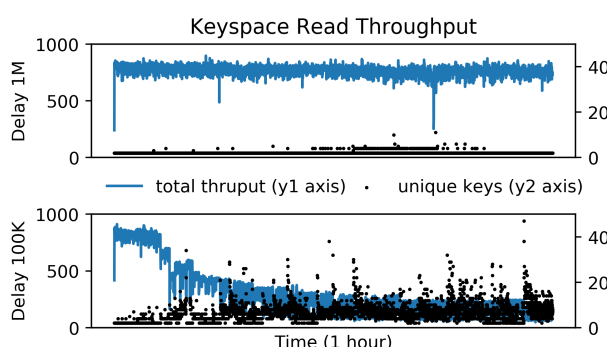


FIG. 1. A load balancing policy that replicates the 10 most recently accessed keys is sufficient for Delay 1M but a more complicated keyspace like Delay 100K needs dynamic load balancing policies.

## 1. Introduction

- key-value stores scale because they support
  1. fine scale annotation
  2. flexible, extensible formats
- science apps are structured, entropy increases over time (e.g., Figure 1 shows key distribution and key popularity changing over time)

Hypothesis: re-distributing keys requires dynamic load balancing policies[1], similar to distributed file systems. A one-size-fits all policy is not sufficient.

## 2. Background

## 3. Methodology

Parsplice is an HPC application with distinct workload phases and a well-known keyspace (Figure 2a).

---

\*Corresponding author address: Los Alamos National Laboratory  
E-mail: msevilla@ucsc.edu

Part 1: Parsplice architecture uses a backend KV store

- Single Node DB (LevelDB, BerkeleyDB) is insufficient
- Distributed KV store solves sync problem and enables load balancing

- HXHIM

Part 2: As Parsplice simulates, entropy increases resulting in keyspace imbalance

- Figure 1 shows how imbalance controlled by “delay”
- Mantle (approach/API) to explore dynamic load balancing policies

- quantifies effect of load balancing
- formalized effective FS balancers
- debugging tool

- HXHIM is a good fit because it has migration mechanisms for load balancing

- bulk operations (put/get())
- key partitioners
- secondary indices

Results should show, In order from most likely to least likely:

1. HPC key-value store workloads are structured (because they are mostly workflows and simulations) that their job phases can be learned and exploited using dynamic load balancing policies.
2. HPC key-value store workloads are so structured that one policy-fits-all
3. HPC key-value store workloads are not structured enough to be learned
4. HPC key-value store workload hotspots/flash crowds are too fast to be exploited




FIG. 2. We can predict how fast the keyspace grows and which parts of the namespace are popular.

#### 4. Conclusion

1. analysis of Parsplice keyspace
2. using a modern distributed kv store
3. positive effects of Mantle

*Acknowledgments.* Start acknowledgments here.

#### References

- [1] D. Perez, E. D. Cubuk, A. Waterland, E. Kaxiras, and A. F. Voter. Long-Time Dynamics Through Parallel Trajectory Splicing. *Journal of chemical theory and computation*.